

Performance Optimization with PerfExpert and MACPO

Jim Browne, Leo Fialho and Ashay Rane

ICS 2013

figures/tacc

figures/ut

Agenda

- ▶ Why another performance tool?
- ▶ What is MACPO?
- ▶ What does MACPO tell you?
- ▶ How to use MACPO?
- ▶ Details of MACPO metrics

figures/tacc

figures/ut

State of the art

- ▶ Modern processors can record performance events
- ▶ Performance events provide fairly accurate view of CPU execution
- ▶ Various tools exist to correlate performance events to user code
- ▶ Examples: PerfExpert, TAU, HPCToolkit, VTune, Scalasca, etc.

figures/tacc

figures/ut

But memory is a bigger problem




Figure 2.2 Starting with 1980 performance as a baseline, the gap in performance, measured as the difference in the time between processor memory requests (for a single processor or core) and the latency of a DRAM access, is plotted over time.

Source: Computer Architecture: A Quantitative Approach, 4th Edition by Hennessy and Patterson

- ▶ For data-intensive applications, application performance usually dependent on memory and not the processor.
- ▶ Performance events provide a CPU-centric, not a memory-centric view.

Performance counter info is limited

- ▶ Performance events are too-fine grained for memory profiling
- ▶ Some performance events are ambiguous to interpret
- ▶ Performance event measurements cannot be scoped to just the important variables

figures/tacc

figures/ut

Fine-grain measurements

- ▶ Performance events are measured on execution of each instruction
- ▶ `mov ah, 16` causes cache miss \Rightarrow increment cache miss counter
- ▶ But memory is optimized for stream traffic and regular accesses (locality, bandwidth, reuse)
- ▶ Hence gap between measurements and optimization techniques

figures/tacc

figures/ut

Ambiguous interpretation

- ▶ Same symptoms but different root causes
- ▶ For instance, L3 cache misses could mean any of the following:
 - ▶ Capacity misses (cold cache)
 - ▶ Poor locality (less reuse of data structures)
 - ▶ False sharing (two processors writing to same cache line)

figures/tacc

figures/ut

Little or no scoping

- ▶ Performance events are triggered for all instructions
- ▶ Hence all memory accesses are profiled
- ▶ But information about only specific data structures is desirable
- ▶ Can greatly speed up problem resolution

figures/tacc

figures/ut


Hence, MACPO

- ▶ Memory Access Centric Performance Optimization
- ▶ Performance tool that analyzes memory access patterns to find sources of inefficiency
- ▶ Used as part of the compilation process
- ▶ Tracks accesses to arrays and structures within a function
- ▶ Tags each such access with source code location and other data
- ▶ Analyzes accesses to see if access patterns can be improved
- ▶ Works with C, C++ and Fortran code [+ Pthreads, OpenMP, MPI]

figures/tacc

figures/ut

MACPO workflow



Combines information from compiler, architecture and from simulation

figures/tacc

figures/ut

What does MACPO tell you?

For each important variable, MACPO shows:

- ▶ Access strides and the frequency of occurrence
- ▶ Presence or absence of cache thrashing and the frequency
- ▶ NUMA misses
- ▶ Reuse factors for data caches

figures/tacc

figures/ut

Sample output

```
Var "counts", seen 1668 times, estimated to cost 147.12 cycles  
Stride of 0 cache lines was observed 1585 times (100.00%).
```

```
Level 1 data cache conflicts = 78.22% [#####]  
Level 2 data cache conflicts = 63.37% [#####]  
NUMA data conflicts = 43.56%           [#####]
```

```
Level 1 data cache reuse factor = 97.0% [#####]  
Level 2 data cache reuse factor = 3.0% [##  
Level 3 data cache reuse factor = 0.0% [
```

figures/tacc

figures/ut

How to use MACPO?

- ▶ Compile application using `macpo.sh`
- ▶ Run application as usual
- ▶ Analyze macpo logs using `macpo-analyze`

figures/tacc

figures/ut

How to use MACPO?

```
# Compile the application using macpo.sh and either --macpo
macpo.sh --macpo:function=thread_func -c mcpi.cc
macpo.sh --macpo:function=thread_func -o mcpi mcpi.o

# Run the application as usual
./mcpi

# Post-process logs to get analysis output
macpo-analyze macpo.out
```

figures/tacc

figures/ut

Understanding MACPO metrics

- ▶ Access strides
- ▶ Cache conflicts
- ▶ NUMA misses
- ▶ Reuse factor for data caches

figures/tacc

figures/ut

Metric #1: Cycles per access

Example:

Var "counts", seen 1073 times, estimated to cost 8.98 cycles

- ▶ Provides estimate of performance impact of accesses to variable
- ▶ Can be used to rule-out variables from further consideration

figures/tacc

figures/ut

Metric #2: Access strides

Example:

Stride of 0 cache lines was observed 983 times (97.62%).

Stride of 2 cache lines was observed 24 times (2.38%).

- ▶ Programs that have unit strides or small regular stride values generally execute fast
- ▶ If stride value is high, look for inverted loops affecting the row-major or column-major ordering

figures/tacc

figures/ut

Metric #3: Cache conflicts

Example:

Level 1 data cache conflicts = 78.22% [#####]

Level 2 data cache conflicts = 63.37% [#####]

- ▶ Indicates multiple cores writing to the same cache line
- ▶ Add dummy bytes to the array so that each processor writes to a different cache line

figures/tacc

figures/ut

Metric #4: NUMA misses

Example:

NUMA data conflicts = 43.56%

[#####]

- ▶ NUMA misses generally arise from one processor initializing all of the shared memory
- ▶ To eliminate NUMA misses, have each processor initialize it's portion of shared memory

figures/tacc

figures/ut

Metric #5: Reuse factors

Example:

Level 1 data cache reuse factor = 94.1% [#####]

Level 2 data cache reuse factor = 5.9% [###

Level 3 data cache reuse factor = 0.0% [

- ▶ Reuse factor indicates the number of times a cache was reused before it was evicted
- ▶ Improve reuse factors by using techniques to improve locality

figures/tacc

figures/ut

Summary

- ▶ MACPO is a tool to analyze memory access patterns
- ▶ NOT a replacement for PerfExpert. Instead, complements PerfExpert's diagnosis.
- ▶ Allows collection of memory traces for arrays and structures
- ▶ Analyzes traces offline to calculate performance metrics
- ▶ This is an early release, so help us squash the bugs! :)

figures/tacc

figures/ut

Sample application

- ▶ Monte-Carlo computation of π
- ▶ Source code online at:
`cs.utexas.edu/users/ashay/xsede-2013/sample-programs.ta`
- ▶ Uses basic C++, parallelized using Pthreads
- ▶ Tasks performed by each thread:
 - ▶ Generates a buffer of random numbers
 - ▶ For each pair of random numbers, calculates z
 - ▶ Checks a condition on z , based on the result increments a counter

figures/tacc

figures/ut

Thread function

```
float x, y, z;
thread_info_t* thread_info = (thread_info_t*) arg;
for (int repeat=0; repeat<REPEAT_COUNT; repeat++)
{
    for (int i=0; i<ITERATIONS; i++)
    {
        x = random_numbers[(i+thread_info->tid)%RANDOM_BUFFER_SIZE];
        y = random_numbers[(1+i+thread_info->tid)%RANDOM_BUFFER_SIZE];

        z = x*x + y*y;
        if (z < 1) counts[thread_info->tid]++;
    }
}
```

MACPO commands for sample code

```
# Compile the application using macpo.sh and either --macpo or
macpo.sh --macpo:function=thread_func source.c compute.c -o mm

# Run the application as usual
./mm

# Post-process logs to get analysis output
macpo-analyze macpo.out
```

figures/tacc

figures/ut

MACPO analysis output

```
macpo-analyze macpo.out
```

```
Var "counts", seen 1668 times, estimated to cost 147.12 cycles  
Stride of 0 cache lines was observed 1585 times (100.00%).
```

```
Level 1 data cache conflicts = 78.22% [#####  
Level 2 data cache conflicts = 63.37% [#####  
NUMA data conflicts = 43.56%           [#####
```

```
Level 1 data cache reuse factor = 97.0% [#####  
Level 2 data cache reuse factor = 3.0% [##  
Level 3 data cache reuse factor = 0.0% [
```

figures/tacc

figures/ut

Problem resolution

- ▶ MACPO shows cache thrashing for `counts` variable.
- ▶ Solution: Add dummy bytes, thus all processors write to different cache lines
- ▶ Optimized code in `monte-carlo-v2.cc`

figures/tacc

figures/ut

MACPO analysis output for optimized code

```
macpo-analyze macpo.out
```

```
Var "counts", seen 1073 times, estimated to cost 8.98 cycles  
Stride of 0 cache lines was observed 983 times (97.62%).  
Stride of 2 cache lines was observed 24 times (2.38%).
```

```
Level 1 data cache conflicts = 0.00% [  
Level 2 data cache conflicts = 0.00% [  
NUMA data conflicts = 0.00%          [
```

```
Level 1 data cache reuse factor = 94.1% [#####  
Level 2 data cache reuse factor = 5.9% [###  
Level 3 data cache reuse factor = 0.0% [
```

figures/tacc

figures/ut

Review

- ▶ Compiled application using `macpo.sh`
- ▶ Discovered cache thrashing for the `counts` array
- ▶ Padding the array reduced cache conflicts from 70% to 0%
- ▶ Execution time improved from 9.14s to 3.17s (65% improvement)

figures/tacc

figures/ut

Summary of MACPO instructions

Compile the application using `macpo.sh` and either `--macpo:1` or `--macpo:2`
`macpo.sh --macpo:function=thread_func source.c compute.c -o compute.o`

Run the application as usual

`./mm`

Post-process logs to get analysis output

`macpo-analyze macpo.out`

figures/tacc

figures/ut