

Performance Optimization for Stampede

Jim Browne, Leo Fialho and Ashay Rane

XSEDE 2013



THE UNIVERSITY OF
TEXAS
AT AUSTIN

Agenda

- ▶ Why another performance tool?
- ▶ What is MACPO?
- ▶ What does MACPO tell you?
- ▶ How to use MACPO?
- ▶ Details of MACPO metrics

State of the art

- ▶ Modern processors can record performance events
- ▶ Performance events provide accurate view of CPU execution
- ▶ Various tools exist to correlate performance events to user code
- ▶ Ex: PerfExpert, TAU, HPCToolkit, VTune, Scalasca, etc.

But memory is a bigger problem than CPU




Figure 2.2 Starting with 1980 performance as a baseline, the gap in performance, measured as the difference in the time between processor memory requests (for a single processor or core) and the latency of a DRAM access, is plotted over time.

Source: Computer Architecture: A Quantitative Approach, 4th Edition by Hennessy and Patterson

But memory is a bigger problem than CPU




Figure 2.2 Starting with 1980 performance as a baseline, the gap in performance, measured as the difference in the time between processor memory requests (for a single processor or core) and the latency of a DRAM access, is plotted over time.

Source: Computer Architecture: A Quantitative Approach, 4th Edition by Hennessy and Patterson

- ▶ For data-intensive applications, application performance usually dependent on memory and not the processor.
- ▶ Performance events provide a CPU-centric, not a memory-centric view.

Limitations for performance events:

#1: Fine-grain measurements

- ▶ Performance events are measured on execution of each instruction
- ▶ `mov ah, 16` causes cache miss \Rightarrow increment cache miss counter
- ▶ But memory is optimized for stream traffic and regular accesses (locality, bandwidth, reuse)
- ▶ Hence gap between measurements and optimization techniques

Limitations for performance events:

#2: Ambiguous interpretation

- ▶ Same symptoms but different root causes
- ▶ For instance, L3 cache misses could mean any of the following:
 - ▶ Capacity misses (cold cache)
 - ▶ Poor locality (less reuse of data structures)
 - ▶ False sharing (two processors writing to same cache line)

Limitations for performance events:

#3: Little or no scoping

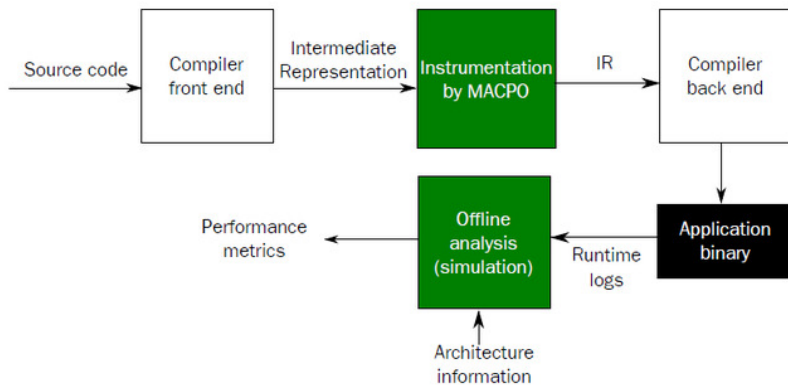
- ▶ Performance events are triggered for all instructions
- ▶ Hence all memory accesses are profiled
- ▶ But information about only specific data structures is desirable
- ▶ Can greatly speed up problem resolution

Hence, MACPO

Memory Access Centric Performance Optimization

- ▶ Analyzes access patterns to find sources of inefficiency
- ▶ Used as part of the compilation process
- ▶ Tracks accesses to arrays and structures within a function
- ▶ Tags each such access with source code location
- ▶ Analyzes accesses to see if access patterns can be improved
- ▶ Works with C, C++ and Fortran code [+ Pthreads, OpenMP]

MACPO workflow



Combines information from compiler, architecture and simulation

What does MACPO tell you?

For each important variable, MACPO shows:

- ▶ Access strides and the frequency of occurrence
- ▶ Presence or absence of cache thrashing and the frequency
- ▶ NUMA misses
- ▶ Reuse factors for data caches

Sample output

Var "counts", seen 1668 times, estimated to cost 147.12 cycles on every access
Stride of 0 cache lines was observed 1585 times (100.00%).

```
Level 1 data cache conflicts = 78.22% [##### ]
Level 2 data cache conflicts = 63.37% [##### ]
NUMA data conflicts = 43.56% [##### ]

Level 1 data cache reuse factor = 97.0% [##### ]
Level 2 data cache reuse factor = 3.0% [## ]
Level 3 data cache reuse factor = 0.0% [ ]
```

How to use MACPO?

- ▶ Compile application using `macpo.sh`
- ▶ Run application as usual
- ▶ Analyze macpo logs using `macpo-analyze`

How to use MACPO?

```
# Compile the application using macpo.sh
# Specify code section using --macpo:function or --macpo:loop
macpo.sh --macpo:function=thread_func -c mcpi.cc
macpo.sh --macpo:function=thread_func -o mcpi mcpi.o

# Run the application as usual
./mcpi

# Post-process logs to get analysis output
macpo-analyze macpo.out
```

Understanding MACPO metrics

- ▶ Access strides
- ▶ Cache conflicts
- ▶ NUMA misses
- ▶ Reuse factor for data caches

Metric #1: Cycles per access

Example:

Var "counts", seen 1073 times,
estimated to cost 8.98 cycles on every access

- ▶ Provides estimate of performance impact of accesses to variable
- ▶ Can be used to rule-out variables from further consideration

Metric #2: Access strides

Example:

Stride of 0 cache lines was observed 983 times (97.62%).

Stride of 2 cache lines was observed 24 times (2.38%).

- ▶ Programs that have unit strides or small regular stride values generally execute fast
- ▶ If stride value is high, look for inverted loops affecting the row-major or column-major ordering

Metric #3: Cache conflicts

Example:

Level 1 data cache conflicts = 78.22%

Level 2 data cache conflicts = 63.37%

- ▶ Indicates multiple cores writing to the same cache line
- ▶ Add dummy bytes to the array so that each processor writes to a different cache line

Metric #4: NUMA misses

Example:

NUMA data conflicts = 43.56%

- ▶ NUMA misses generally arise from one processor initializing all of the shared memory
- ▶ To eliminate NUMA misses, have each processor initialize it's portion of shared memory

Metric #5: Reuse factors

Example:

Level 1 data cache reuse factor = 94.1%

Level 2 data cache reuse factor = 5.9%

Level 3 data cache reuse factor = 0.0%

- ▶ Reuse factor indicates the number of times a cache was reused before it was evicted
- ▶ Improve reuse factors by using techniques to improve locality

Summary

- ▶ MACPO is a tool to analyze memory access patterns
- ▶ NOT a replacement for PerfExpert. Instead, complements PerfExpert's diagnosis.
- ▶ Allows collection of memory traces for arrays and structures
- ▶ Analyzes traces offline to calculate performance metrics
- ▶ This is an early release, so help us squash the bugs! :)

Sample application

- ▶ Monte-Carlo computation of Pi
- ▶ Source code online at: <http://goo.gl/uEVrh>
- ▶ Uses basic C++, parallelized using Pthreads
- ▶ Tasks performed by each thread:
 - ▶ Generates a buffer of random numbers
 - ▶ For each pair of random numbers, calculates z
 - ▶ Checks a condition on z , based on the result increments a counter

Thread function

```
int t;
float x, y, z;
thread_info_t* thread_info = (thread_info_t*) arg;
for (int repeat=0; repeat<REPEAT_COUNT; repeat++)
{
    for (int i=0; i<ITERATIONS; i++)
    {
        t = i+thread_info->tid;
        x = random_numbers[t%RANDOM_BUFFER_SIZE];
        y = random_numbers[(1+t)%RANDOM_BUFFER_SIZE];

        z = x*x + y*y;
        if (z < 1) counts[thread_info->tid]++;
    }
}
```

MACPO commands for sample code

```
# Compile the application using macpo.sh
# Specify code section using --macpo:function or --macpo:loop
macpo.sh --macpo:function=thread_func source.c compute.c -o mm

# Run the application as usual
./mm

# Post-process logs to get analysis output
macpo-analyze macpo.out
```


MACPO analysis output (truncated)

```
macpo-analyze macpo.out
```

```
Var "counts", seen 1668 times,  
estimated to cost 147.12 cycles on every access
```

```
Stride of 0 cache lines was observed 1585 times (100.00%).
```

```
Level 1 data cache conflicts = 78.22%
```

```
Level 2 data cache conflicts = 63.37%
```

```
NUMA data conflicts = 43.56%
```

```
Level 1 data cache reuse factor = 97.0%
```

```
Level 2 data cache reuse factor = 3.0%
```

```
Level 3 data cache reuse factor = 0.0%
```

Problem resolution

- ▶ MACPO shows cache thrashing for `counts` variable.
- ▶ Solution: Add dummy bytes, thus all processors write to different cache lines
- ▶ Optimized code in `monte-carlo-v2.cc`

MACPO analysis output for optimized code (truncated)

```
macpo-analyze macpo.out
```

```
Var "counts", seen 1073 times,  
estimated to cost 8.98 cycles on every access
```

```
Stride of 0 cache lines was observed 983 times (97.62%).  
Stride of 2 cache lines was observed 24 times (2.38%).
```

```
Level 1 data cache conflicts = 0.00%  
Level 2 data cache conflicts = 0.00%  
NUMA data conflicts = 0.00%
```

```
Level 1 data cache reuse factor = 94.1%  
Level 2 data cache reuse factor = 5.9%  
Level 3 data cache reuse factor = 0.0%
```

Review

- ▶ Compiled application using `macpo.sh`
- ▶ Discovered cache thrashing for the `counts` array
- ▶ Padding the array reduced cache conflicts from 70% to 0%
- ▶ Execution time: 9.14s to 3.17s (65% improvement)

Summary of MACPO instructions

```
# Compile the application using macpo.sh
# Specify code section using --macpo:function or --macpo:loop
macpo.sh --macpo:function=thread_func source.c compute.c -o mm

# Run the application as usual
./mm

# Post-process logs to get analysis output
macpo-analyze macpo.out
```